

CREDIT RISK PREDICTION

Problem Statement:

The goal is to build a Credit Risk Prediction Model that determines whether a borrower will default (1) or not default (0) on a loan. This helps lenders mitigate financial risks.

```
In [1]: import pandas as pd # for data manipulation
import numpy as np # for Mathematical Operations
```

```
In [2]: path = 'C:/Users/Sanayak/Desktop/Credit Risk Dataset/credit_risk_dataset.csv'
df_loan = pd.read_csv(path)
df_loan.sample(10)
```

```
Out[2]:
```

	person_age	person_income	person_home_ownership	person_emp_length	loan_intent	loan_grade	loan_amnt
27289	27	130000	MORTGAGE	1.0	HOMEIMPROVEMENT	B	21000
8720	25	60000	MORTGAGE	6.0	HOMEIMPROVEMENT	B	4000
1844	21	30000	MORTGAGE	0.0	MEDICAL	A	7125
10305	22	66396	MORTGAGE	6.0	EDUCATION	C	15000
32564	53	45000	RENT	0.0	PERSONAL	C	5600
19527	32	49000	RENT	2.0	HOMEIMPROVEMENT	B	12000
14410	23	105600	MORTGAGE	2.0	PERSONAL	A	25000
6423	22	40000	RENT	1.0	MEDICAL	B	6000
28555	27	80000	MORTGAGE	7.0	MEDICAL	B	7400
18573	33	67800	RENT	3.0	HOMEIMPROVEMENT	B	18000

Variables Definition:

1. Person_age (Customer's age): Helps identify age groups that may be more prone to default.

2. Person_income (Annual income): The client's repayment capacity is strongly associated with their income.
3. Person_home_ownership (Home ownership): Indicates whether the client owns a home, rents, is mortgaged, etc. This influences financial stability.
4. Person_emp_length (Years of employment): Longer employment tenure often corresponds to lower default risk.
5. Loan_intent (Loan intent): The purpose of the loan (e.g., education, car, home) can influence the risk, as different loan types have varying default probabilities.
6. Loan_grade (Loan grade): An internal rating assigned to the loan indicating the initial perceived risk level.
7. Loan_amnt (Loan amount): Higher loan amounts mean greater exposure to risk.
8. Loan_int_rate (Interest rate): Higher interest rates are often associated with higher default risk.
9. loan_status (Loan status – target variable): 0 means non-default, 1 means default.
10. Loan_percent_income (Percent of income dedicated to the loan): A metric relating the loan amount to the borrower's income, indicating financial burden.
11. Cb_person_default_on_file (Historical defaults): Indicates if the client has previously defaulted.
12. Cb_person_cred_hist_length (Length of credit history): A longer credit history provides better insight into a client's behavior.

```
In [3]: print('Number of Rows by Columns:',df_loan.shape)# Get the shape of the dataset
print('Number of duplicates:',df_loan.duplicated().sum())# Check for duplicated values
```

Number of Rows by Columns: (32581, 12)

Number of duplicates: 165

```
In [4]: df_loan.drop_duplicates(inplace=True)#drop the duplicated entries and update the dataset
print(df_loan.duplicated().sum())
```

0

```
In [5]: print('The number of missing values for each column are: ')
print(df_loan.isna().sum())#Check for null values
```

The number of missing values for each column are:

```

person_age          0
person_income       0
person_home_ownership 0
person_emp_length   887
loan_intent         0
loan_grade         0
loan_amnt          0
loan_int_rate      3095
loan_status        0
loan_percent_income 0
cb_person_default_on_file 0
cb_person_cred_hist_length 0
dtype: int64

```

In [6]: `df_loan.describe()`

Out[6]:

	person_age	person_income	person_emp_length	loan_amnt	loan_int_rate	loan_status	loan_percent_income	cb_pe
count	32416.000000	3.241600e+04	31529.00000	32416.000000	29321.000000	32416.000000	32416.000000	
mean	27.747008	6.609164e+04	4.79051	9593.845632	11.017265	0.218688	0.170250	
std	6.354100	6.201558e+04	4.14549	6322.730241	3.241680	0.413363	0.106812	
min	20.000000	4.000000e+03	0.00000	500.000000	5.420000	0.000000	0.000000	
25%	23.000000	3.854200e+04	2.00000	5000.000000	7.900000	0.000000	0.090000	
50%	26.000000	5.500000e+04	4.00000	8000.000000	10.990000	0.000000	0.150000	
75%	30.000000	7.921800e+04	7.00000	12250.000000	13.470000	0.000000	0.230000	
max	144.000000	6.000000e+06	123.00000	35000.000000	23.220000	1.000000	0.830000	

In [7]: `# Fill null with median value for the individual columns`
`df_loan[['person_emp_length', 'loan_int_rate']] = df_loan[['person_emp_length',`
 `'loan_int_rate']].fillna(df_loan[['person_emp_length',`
 `'loan_int_rate']].median()`

In [8]: `df_loan[['person_emp_length', 'loan_int_rate']].isna().sum()`

```
Out[8]: person_emp_length    0
        loan_int_rate        0
        dtype: int64
```

```
In [9]: df_loan.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 32416 entries, 0 to 32580
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   person_age                            32416 non-null  int64
1   person_income                         32416 non-null  int64
2   person_home_ownership                 32416 non-null  object
3   person_emp_length                     32416 non-null  float64
4   loan_intent                           32416 non-null  object
5   loan_grade                           32416 non-null  object
6   loan_amnt                            32416 non-null  int64
7   loan_int_rate                         32416 non-null  float64
8   loan_status                          32416 non-null  int64
9   loan_percent_income                  32416 non-null  float64
10  cb_person_default_on_file             32416 non-null  object
11  cb_person_cred_hist_length            32416 non-null  int64
dtypes: float64(3), int64(5), object(4)
memory usage: 3.2+ MB
```

Observations:

The Age column has an Highest value of 144 and the employment length has an Highest value of 123 which might be caused due to human error

```
In [10]: # Define a function that removes the 95 percentile
def remove_outliers(df, columns):
    for col in columns:
        threshold = df[col].quantile(0.95)
        df = df[df[col] <= threshold]
    return df

cleaned_df = remove_outliers(df_loan, columns=['person_age', 'person_emp_length'])
print('Original DataFrame:')
print(df_loan[['person_age', 'person_emp_length']].max())
```

```
print('Cleaned DataFrame:')  
print(cleaned_df[['person_age', 'person_emp_length']].max())
```

Original DataFrame:

```
person_age      144.0  
person_emp_length 123.0  
dtype: float64
```

Cleaned DataFrame:

```
person_age      40.0  
person_emp_length 12.0  
dtype: float64
```

Outliers:

Outliers were identified in:

`person_age` : Maximum = **144** (unrealistic).

`person_emp_length` : Maximum = **123**.

To fix this, values above the 95th percentile were removed.

Post-cleaning:

`person_age` capped at **40**.

`person_emp_length` capped at **12**.

```
In [11]: cleaned_df.describe()
```


Out[11]:

	person_age	person_income	person_emp_length	loan_amnt	loan_int_rate	loan_status	loan_percent_income	cb_pe
count	29543.000000	2.954300e+04	29543.000000	29543.000000	29543.000000	29543.000000	29543.000000	
mean	26.557729	6.401581e+04	4.207325	9497.601801	11.028543	0.222015	0.171306	
std	4.410420	4.431807e+04	3.123315	6262.114024	3.076598	0.415608	0.106809	
min	20.000000	4.080000e+03	0.000000	500.000000	5.420000	0.000000	0.000000	
25%	23.000000	3.800000e+04	2.000000	5000.000000	8.490000	0.000000	0.090000	
50%	25.000000	5.500000e+04	4.000000	8000.000000	10.990000	0.000000	0.150000	
75%	29.000000	7.800000e+04	6.000000	12000.000000	13.160000	0.000000	0.230000	
max	40.000000	1.200000e+06	12.000000	35000.000000	23.220000	1.000000	0.830000	


Exploratory Data Analysis

```
In [12]: import matplotlib.pyplot as plt
import seaborn as sns
```


```
In [13]: plt.figure(figsize=(6,4))
ax = sns.countplot(x='loan_status', data=cleaned_df, hue='loan_status')
for p in ax.patches:
    ax.annotate(f'{p.get_height()}',
                (p.get_x() + p.get_width() / 2., p.get_height()), ha='center', va='bottom')
plt.title('Loan Status Distribution')
plt.xlabel('Loan Status')
plt.ylabel('Number of Loans')
plt.xticks(ticks=[0,1], labels=['Non-Default', 'Default'])
plt.legend()
plt.show()
```




```
In [14]: plt.figure(figsize=(6,4))
sns.histplot(cleaned_df['person_age'],kde=False,bins=20)
plt.title('Age Distribution')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.show()
```



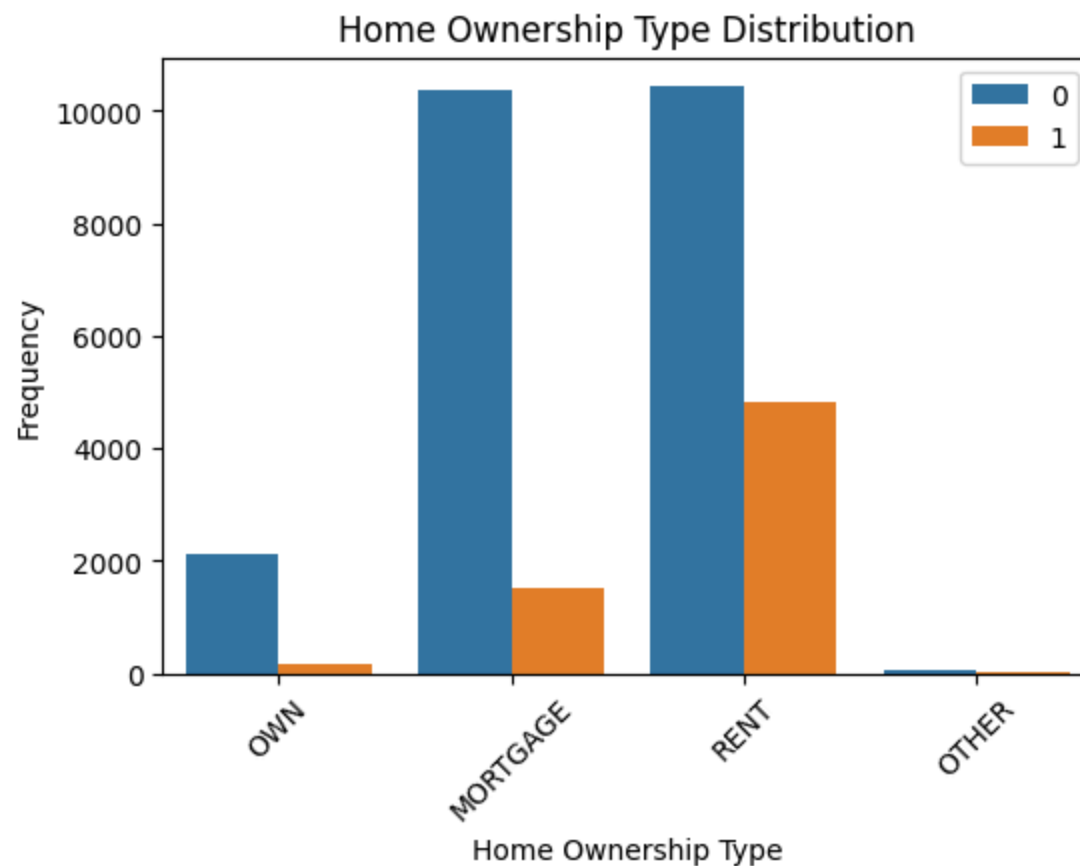
```
In [15]: loan_intent_counts=cleaned_df['loan_intent'].value_counts()
labels = ['PERSONAL', 'VENTURE', 'EDUCATION', 'MEDICAL', 'HOMEIMPROVEMENT', 'DEBTCONSOLIDATION']
plt.figure(figsize=(6,4))
plt.pie(loan_intent_counts, labels=labels, autopct="%1.1f%%")
plt.title('Sectorial Representation of Loan Intent')
plt.show()
```

```
In [16]: plt.figure(figsize=(6,4))
sns.countplot(x='loan_intent', data=cleaned_df,hue='loan_status')
plt.title('Loan Intent Distribution')
plt.xlabel('Loan Intent')
plt.ylabel('Frequency')
plt.xticks(rotation=45)
plt.legend()
plt.show()
```



```
In [17]: plt.figure(figsize=(6,4))
sns.countplot(x='person_home_ownership', data=cleaned_df, hue='loan_status')
plt.title('Home Ownership Type Distribution')
plt.xlabel('Home Ownership Type')
plt.ylabel('Frequency')
plt.xticks(rotation=45)
plt.legend()
plt.show()
```



```
In [18]: from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.pipeline import Pipeline
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
```

```
In [19]: # Split in to feature(X) and target variable(y)
X = cleaned_df.drop('loan_status', axis=1)
y = cleaned_df['loan_status']

# Split in to training and testing set
X_train, X_test, y_train, y_test = train_test_split(X, y,
```

```
test_size = 0.25,  
random_state = 42)
```

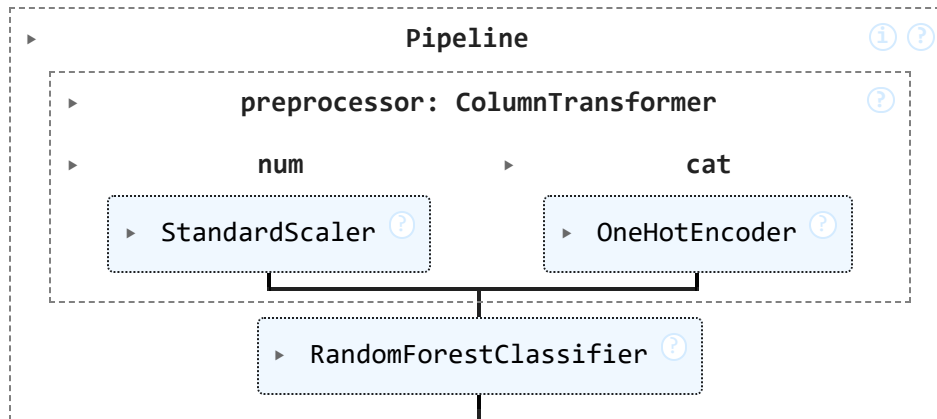
```
In [20]: # Identify the categorical and numerical features  
categorical_features = X.select_dtypes(include=['object']).columns  
numerical_features = X.select_dtypes(include=['int64', 'float64']).columns  
  
# Create preprocessing pipeline for categorical and numerical features  
categorical_transform = Pipeline(steps=[('ohe', OneHotEncoder(handle_unknown='ignore'))])  
numerical_transform = Pipeline(steps=[('scaler', StandardScaler())])  
  
#Combine the processing pipelines  
preprocessor = ColumnTransformer(transformers=[('num', numerical_transform,  
                                              numerical_features), ('cat', categorical_transform,  
                                                                    categorical_features)])
```

Data Preprocessing

- **Categorical Features:**
 - Encoded using **OneHotEncoder**.
- **Numerical Features:**
 - Standardized using **StandardScaler**.
- Split data into **train (75%)** and **test (25%)** sets.

```
In [21]: # Create a pipeline for the xgboost classifier  
xgb_pipeline = Pipeline(steps=[('preprocessor', preprocessor), ('classifier', XGBClassifier(random_state=42))])  
  
# Create a pipeline for the randomforest classifier  
rf_pipeline = Pipeline(steps=[('preprocessor', preprocessor), ('classifier', RandomForestClassifier(random_state=42))])  
  
#Train the xgboost model  
xgb_pipeline.fit(X_train, y_train)  
  
# Train the randomforest classifier  
rf_pipeline.fit(X_train, y_train)
```

Out[21]:



Model Training

Two classifiers were built using pipelines:

1. **XGBoost.**
2. **Random Forest.**

Both models were trained and tested on the cleaned dataset.

```

In [22]: # Make predictions on the testset
xgb_y_pred = xgb_pipeline.predict(X_test)
rf_y_pred = rf_pipeline.predict(X_test)

xgb_accuracy = accuracy_score(xgb_y_pred, y_test)
rf_accuracy = accuracy_score(rf_y_pred, y_test)

xgb_class_report = classification_report(xgb_y_pred, y_test)
rf_class_report = classification_report(rf_y_pred, y_test)

In [23]: print(f'XGBoost Model Accuracy: {xgb_accuracy:.2f}')
print(f'RandomForest Model Accuracy: {rf_accuracy:.2f}')

print('XGBoost Classification Report:\n', xgb_class_report)
print('RandomForest Classification Report:\n', rf_class_report)

```

XGBoost Model Accuary: 0.93

RandomForest Model Accuary: 0.93

XGBoost Classification Report:

	precision	recall	f1-score	support
0	0.99	0.93	0.96	6059
1	0.75	0.95	0.84	1327
accuracy			0.93	7386
macro avg	0.87	0.94	0.90	7386
weighted avg	0.95	0.93	0.94	7386

RandomForest Classification Report:

	precision	recall	f1-score	support
0	0.99	0.92	0.96	6131
1	0.72	0.97	0.83	1255
accuracy			0.93	7386
macro avg	0.86	0.95	0.89	7386
weighted avg	0.95	0.93	0.94	7386

- Both models achieved **93% accuracy**.
- XGBoost performed slightly better on recall for defaults.

```
In [24]: from sklearn.model_selection import GridSearchCV

# Define parameter grid for XGBoost and Random Forest
param_grid_xgb = {'classifier__max_depth': [3,5,7],
                  'classifier__learning_rate':[0.1,0.01,0.001],
                  'classifier__n_estimators':[100,200,300]}

param_grid_rf = {'classifier__n_estimators':[100,200,300],
                 'classifier__max_depth':[5,10,15],
                 'classifier__min_samples_split':[2,5,10],
                 'classifier__min_samples_leaf':[1,2,4]}

# Create GridSearch objects
grid_search_xgb = GridSearchCV(xgb_pipeline, param_grid_xgb, cv=5, scoring='accuracy')
```

```
grid_search_rf = GridSearchCV(rf_pipeline, param_grid_rf, cv=5, scoring='accuracy')

# Fit the GridSearch object to training data
grid_search_xgb.fit(X_train, y_train)
grid_search_rf.fit(X_train, y_train)

# Print the best parameters and score for each model
print('Best parameters for XGBoost:', grid_search_xgb.best_params_)
print('Best Score for XGBoost:', grid_search_xgb.best_score_)
print('Best parameters for Random Forest:', grid_search_rf.best_params_)
print('Best Score for Random Forest:', grid_search_rf.best_score_)
```

Best parameters for XGBoost: {'classifier__learning_rate': 0.1, 'classifier__max_depth': 7, 'classifier__n_estimators': 200}

Best Score for XGBoost: 0.9346935196478373

Best parameters for Random Forest: {'classifier__max_depth': 15, 'classifier__min_samples_leaf': 1, 'classifier__min_samples_split': 5, 'classifier__n_estimators': 300}

Best Score for Random Forest: 0.9314440046211994

```
In [25]: # use the best model to make predictions and evaluate
best_xgb_model = grid_search_xgb.best_estimator_
best_rf_model = grid_search_rf.best_estimator_

xgb_predictions = best_xgb_model.predict(X_test)
rf_predictions = best_rf_model.predict(X_test)

xgb_class_report_grid = classification_report(xgb_predictions,y_test)
rf_class_report_grid = classification_report(rf_predictions,y_test)

# Print the classification report for both models
print('XGBoost Classification Report:\n', xgb_class_report_grid)
print('RandomForest Classification Report:\n', rf_class_report_grid)
```

XGBoost Classification Report:

	precision	recall	f1-score	support
0	0.99	0.93	0.96	6099
1	0.74	0.97	0.84	1287
accuracy			0.94	7386
macro avg	0.87	0.95	0.90	7386
weighted avg	0.95	0.94	0.94	7386

RandomForest Classification Report:

	precision	recall	f1-score	support
0	0.99	0.92	0.96	6144
1	0.72	0.98	0.83	1242
accuracy			0.93	7386
macro avg	0.86	0.95	0.89	7386
weighted avg	0.95	0.93	0.94	7386

Hyperparameter Tuning

- Performed GridSearchCV to optimize model parameters:
 - XGBoost:**
 - Best Parameters: `max_depth=7` , `learning_rate=0.1` , `n_estimators=200` .
 - Best Accuracy: **93.47%**.
 - Random Forest:**
 - Best Parameters: `max_depth=15` , `min_samples_split=5` , `n_estimators=300` .
 - Best Accuracy: **93.14%**.

Final Results

- After tuning, XGBoost outperformed Random Forest slightly in recall and accuracy.
- XGBoost** is the preferred model for Credit Risk Prediction.

Additional Considerations:

-**Model Selection:** Experiment with different models.

-**Regularization:** Apply techniques like L1 and L2 regularization to prevent overfitting.

Conclusion:

The project successfully built and optimized a **Credit Risk Prediction Model**.

Key takeaways:

- Age and loan intent play significant roles in default likelihood.
- Data cleaning and feature engineering improved model performance.
- XGBoost, with hyperparameter tuning, delivered the best results.

```
In [26]: import joblib
```

```
In [27]: joblib.dump(best_xgb_model, 'best_xgb_model.pkl')
```

```
Out[27]: ['best_xgb_model.pkl']
```

Prepared By: Victor Itinah Iniobong

Tools Used: Python, Pandas, Scikit-Learn, XGBoost, Matplotlib, Seaborn.